SQL - the Dark Side

Tamas Piros

Drawbacks of SQL

- They are not designed to handle change
- They are not designed for heterogeneous data
- They are not designed for scale
- They are note designed for mixed workloads
- They are a mismatch for modern app development

RDBMS are not designed to handle change

- Data is created faster and it also changes faster
- Businesses must keep up with this pace
  - Market dynamics, on-demand services, acquisitions
- Data modelling is a lengthy and daunting process in itself, change(s) add(s) further complexity
- Just remember, data models almost always change

Matt Allen (MarkLogic) - http://www.marklogic.com/blog/relational-databases-change/

RDBMS are not designed for heterogeneous data

- Businesses need to store everything
  - Customers, partners, regulators
- Creation of data silos to store heterogeneous data is not the right solution
- Most relational databases only support certain application or certain type of data
- “duct tape” together data sources which leads to additional complexity and frustration
- Unstructured data is complex by nature

Matt Allen (MarkLogic) - http://www.marklogic.com/blog/relational-databases-heterogeneous-data/

RDBMS are not designed for scale

- Scalability and elasticity is challenging for relational systems
- Relational systems are designed to run on single servers in order to maintain table mapping integrity and to avoid distributed computing
    - Scaling is therefore achieved by buying more expensive hardware
- Database vendors apply improvements such as master-slave model, sharding, shared storage, distributed caching
- Flexibility is often sacrificed for performance

Matt Allen (MarkLogic) - http://www.marklogic.com/blog/relational-databases-scale/

RDBMS are not designed for mixed workloads

- Relational databases are designed to handle two types of workloads: operational and analytical but they can't do both
  - OLTP (online transaction processing) is designed for fast transactions - "who bought X?"
  - OLAP (online analytical processing) is designed for aggregating information – "How many people bought X?"
- Data is moved from an operational system to a central data warehouse, then moved again to a data mart to answer specific questions
- Problems start to creep in – maintaining systems, enforcing security

Matt Allen (MarkLogic) - http://www.marklogic.com/blog/relational-databases-are-not-designed-for-mixed-workloads/

RDBMS are a mismatch for modern app development

- ORMs are required – adds mor
  process
- Performance loss, potentially
  difficult to maintain codebase
- Legacy databases require lega

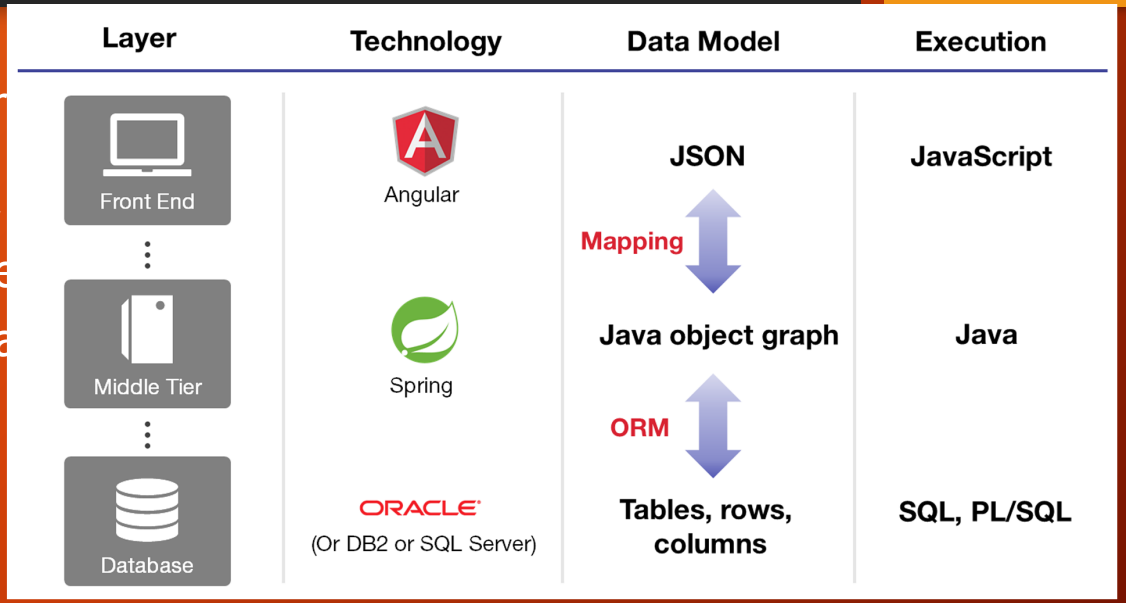

Matt Allen (MarkLogic) - http://www.marklogic.com/blog/rdbms-mismatch-modern-app-development/

Conclusion

- Relational databases do have their own use cases
  - When data is truly relational use relational databases
- In today's rapidly changing environments, where complex data is used, relational databases lag behind
- Remember: use the right technology to the right problem